

Dockerfiles

Florent Gluck - Florent.Gluck@hesge.ch

April 13, 2022

- Docker can build images automatically by reading instructions from a Dockerfile
- **Dockerfile = “source code” or recipe on how to build an image**
- Does not install an image → instructions on how to build an image

Why a Dockerfile?

- When existing images don't satisfy our needs
- When we can't find the exact image we need
- Most images are generic and won't cover our exact scenario
- To customize an existing image to our needs

1. **Write a Dockerfile** with instructions on how to build the image

Image creation and usage workflow

1. **Write** a **Dockerfile** with instructions on how to build the image
2. **Build** the **image** from the Dockerfile

Image creation and usage workflow

1. **Write** a **Dockerfile** with instructions on how to build the image
2. **Build** the **image** from the Dockerfile
3. **Create** the **container** from the newly built image

1. Define base image (**FROM**)

Dockerfile contents workflow

1. Define base image (**FROM**)
2. Set environment variables (**ENV**)

Dockerfile contents workflow

1. Define base image (**FROM**)
2. Set environment variables (**ENV**)
3. Add instructions, e.g. packages to install, etc. (**RUN**)

Dockerfile contents workflow

1. Define base image (**FROM**)
2. Set environment variables (**ENV**)
3. Add instructions, e.g. packages to install, etc. (**RUN**)
4. Add files (**COPY** or **ADD**)

Dockerfile contents workflow

1. Define base image (**FROM**)
2. Set environment variables (**ENV**)
3. Add instructions, e.g. packages to install, etc. (**RUN**)
4. Add files (**COPY** or **ADD**)
5. Define default command (**CMD** or **ENTRYPOINT**)

Example of Dockerfile

```
FROM alpine:3.12
ENV EDITOR=vi
RUN apk update
RUN apk add git
COPY hosts /etc/
WORKDIR /home
CMD git
```

Main commands

- **FROM**: base image to use
- **ENV**: define environment variables (in Dockerfile and container exec.)
- **RUN**: execute commands
- **COPY** / **ADD**: add files to the image
- **CMD** / **ENTRYPOINT**: command to execute when instantiating the image
- **WORKDIR**: working directory (in Dockerfile and container exec.), / by default
- **EXPOSE**: indicates ports the container listens to for connections

COPY vs ADD

- **COPY** / **ADD** both copy files from a specific location into a Docker image
- **COPY** can only copy a local file or directory from the host into the image
- **ADD** similar to **COPY** but more powerful:
 - can extract a **tar** archive into the image
 - can specify an URL instead of a local file or directory

CMD

- **CMD** specify which command to execute when the image is instantiated Preferred form:

```
CMD ["executable", "arg1", "arg2", ...]
```

Shell form, which executes **command** with **/bin/sh -c**:

```
CMD command param1 param2
```

- **CMD can be overridden** when starting the container:

```
docker run [OPTIONS] IMAGE[:TAG] [COMMAND] [ARG...]
```

- **Only** the last **CMD** of a Dockerfile is executed!

ENTRYPOINT

- **ENTRYPOINT** also specify which command to execute when the image is instantiated
- By opposition to **CMD**, **ENTRYPOINT** **cannot be overridden** when starting the container!
- What is the purpose of **ENTRYPOINT** since **CMD** already exists?
 - **CMD** contents is added to **ENTRYPOINT** as argument
 - **ENTRYPOINT** + **CMD** = a way of specifying **default**, but **overridable** arguments

```
ENTRYPOINT ["git"]  
CMD ["--help"]
```


Examples (1/2)

```
FROM alpine
ENTRYPOINT ["/bin/echo"]
```

- `docker run myecho` → (nothing printed)
- `docker run myecho blah` → `blah`
- `docker run myecho blah blah` → `blah blah`

```
FROM alpine
CMD ["/bin/echo"]
```

- `docker run myecho` → (nothing printed)
- `docker run myecho blah` → `error: executable "blah" not found!`
- `docker run myecho /bin/echo blah` → `blah`

Examples (2/2)

```
FROM alpine
ENTRYPOINT ["/bin/echo"]
CMD ["blah"]
```

- `docker run myecho` → `blah`
- `docker run myecho pipo molo` → `pipo molo`
- `docker run myecho blah` → `blah`

```
FROM alpine
```

- `docker run myecho /bin/echo` → (nothing printed)
- `docker run myecho /bin/echo blah` → `blah`

Building an image

- The `docker build` command builds an image from a Dockerfile and a build context
- **The build is run by the Docker daemon, not the client!**
- Example:

```
docker build . -t myimage:beta
```

Build context

- Build context = set of files at a specified PATH (local directory) or URL (git repository)
- A context is processed recursively:
 - PATH includes any subdirectories
 - URL includes a repository and its submodules
- Example: specify `/tmp/pipo/` as the build context:

```
docker build /tmp/pipo
```

- IMPORTANT: during build, the client **sends** the **entire context recursively** to the Docker daemon!
- Use `.dockerignore` to specify which files must not be sent to the daemon (similarly to `.gitignore`)

Building an image from a local root filesystem

- Using the **FROM** directive to specify a base image might not always be desirable
- The image will likely **change** over time
 - a specific tag doesn't prevent the image from being updated
 - bugs might be introduced
 - behavior might slightly change, enough to break things
- Alternative: instead of using a base image from a repository, use an immutable local archive as root filesystem

Exporting an image or container's filesystem

- `docker export` → exports a **container's** filesystem into a tar archive (to `stdout`)
 - archive contains the filesystem only
- `docker save` → saves an **image's** filesystem into a tar archive (to `stdout`)
 - archive contains the various layers' contents (filesystem) and image meta-data (e.g. entry-point, etc.)
- Note that `docker import` and `docker load` perform the opposite operations

Docker export: archive's contents

```
$ docker run --name myc openjdk:latest
$ docker export myc > rootfs.tar && tar tf rootfs.tar
.dockerenv
bin
boot/
dev/
dev/console
dev/full
dev/initctl
dev/null
dev/ptmx
dev/pts/
dev/random
dev/shm/
dev/tty
dev/tty0
dev/urandom
dev/zero
etc/
etc/aliases
etc/alternatives/
etc/alternatives/jar
etc/alternatives/jarsigner
etc/alternatives/java
etc/alternatives/javac
...
```

Docker save: archive's contents

```
$ docker save openjdk:latest > rootfs.tar && tar tf rootfs.tar
34aba91dbd1358ac48d86995dad4620c73ead6466f94f8dfce622a59892fcb5f.json
8e3b009939a813b63c7c2bae06327fa868cdacb2f33edf524d436a1be3036b9a/
8e3b009939a813b63c7c2bae06327fa868cdacb2f33edf524d436a1be3036b9a/VERSION
8e3b009939a813b63c7c2bae06327fa868cdacb2f33edf524d436a1be3036b9a/json
8e3b009939a813b63c7c2bae06327fa868cdacb2f33edf524d436a1be3036b9a/layer.tar
b04eff89da618fd519087acde0f769f144c30e8b3b6c21cf2310248d24c52015/
b04eff89da618fd519087acde0f769f144c30e8b3b6c21cf2310248d24c52015/VERSION
b04eff89da618fd519087acde0f769f144c30e8b3b6c21cf2310248d24c52015/json
b04eff89da618fd519087acde0f769f144c30e8b3b6c21cf2310248d24c52015/layer.tar
eeffb4b5e0bcf55f75dfdc77f8c0c5e4cfaf98e8ff48d350c9ac75768cd019631/
eeffb4b5e0bcf55f75dfdc77f8c0c5e4cfaf98e8ff48d350c9ac75768cd019631/VERSION
eeffb4b5e0bcf55f75dfdc77f8c0c5e4cfaf98e8ff48d350c9ac75768cd019631/json
eeffb4b5e0bcf55f75dfdc77f8c0c5e4cfaf98e8ff48d350c9ac75768cd019631/layer.tar
manifest.json
repositories
```


Creating an image from a local root filesystem

- A local root filesystem archive can be used as base image
- Requires the use of a special empty image named **scratch**
- Steps to create an image from a local root filesystem:
 - 1) Export a container's filesystem into a **tar** archive

```
docker export alpine:3.12 > alpine_3.12.tar
```

- 2) Create a Dockerfile that uses the archive as the root filesystem

```
FROM scratch
ADD alpine_3.12.tar /
CMD sh
```

- 3) Build the image from the Dockerfile

Multi-stage builds

- What if you want to build a container with a specific program that must be generated from source?
- Resulting image would be very large → requires the full build environment!
- How to make the resulting image as small as possible?
 - unfortunately, no easy, clean and generic way of doing so. . .
- Solution ?
 - multi-stage builds!

Multi-stage builds: why?

Purpose of multi-stage builds:

- When images require building binaries or artifacts
- Help keep image size minimal

Multi-stage builds: how?

```
# Builder stage
FROM golang AS builder
ADD . /app
WORKDIR /app
RUN go build

# Final stage which uses the builder stage
FROM alpine:latest
RUN apk --no-cache add ca-certificates
WORKDIR /root/
COPY --from=builder /app/app .
CMD ["/app"]
```

Best practices

- One container should only solve one problem!
- Minimize number of layers (= minimize number of steps)
- Use `.dockerignore` to avoid sending all context files/dirs to the Docker daemon
- Create Dockerfiles that define stateless images
 - Any state should be kept outside of the container
- Order layers from less frequently changed to more frequently changed (ensure build cache is reusable)

Resources

- Dockerfile reference (official)
<https://docs.docker.com/engine/reference/builder/>
- Dockerfile tutorial by example - basics and best practices
<https://takacsmark.com/dockerfile-tutorial-by-example-dockerfile-best-practices-2018/>
- Best practices for writing Dockerfiles (official)
https://docs.docker.com/develop/develop-images/dockerfile_best-practices/
- Create a base image and multi-stage builds (official)
<https://docs.docker.com/develop/develop-images/baseimages/>
<https://docs.docker.com/develop/develop-images/multistage-build/>
- Explaining Docker Image IDs
<https://windsock.io/explaining-docker-image-ids/>